

1)

Size = $6 - 3 + 1 = 4$ By 4Trainable parameters = $((\text{Size of image}) - (\text{size of filter})) * \text{stride} + 1 = (6*6 - 2*2) * 2 + 1 = 65$

2)

with padding 1 around, it would be a $8 * 8$.size = $8 - 4 + 1 = 5$ by 5Trainable parameters = $((\text{Size of image}) - (\text{size of filter})) * \text{stride} + 1 = (8*8 - 2*2) * 2 + 1 = 121$

3)

3) a)
$$\begin{matrix} \text{filter} \\ \begin{bmatrix} -1 & -1 & -1 \\ 1 & 1 & 1 \\ -1 & -1 & -1 \end{bmatrix} \end{matrix}$$

$$\begin{bmatrix} 2 & 5 & 9 \\ 1 & 2 & 0 \\ 1 & 2 & 2 \end{bmatrix} \begin{bmatrix} -1 & -1 & -1 \\ 1 & 1 & 1 \\ -1 & -1 & -1 \end{bmatrix} = 9$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 5 & 9 & 7 \\ 2 & 0 & 1 \end{bmatrix} \begin{bmatrix} -1 & -1 & -1 \\ 1 & 1 & 1 \\ -1 & -1 & -1 \end{bmatrix} = 12$$

and so on

result
$$\begin{bmatrix} 9 & 12 & 7 & -1 \\ -18 & -23 & -26 & -16 \\ 8 & -3 & -2 & -10 \\ -5 & -8 & -14 & -8 \end{bmatrix}$$
 captures horizontal lines features

b)
$$\begin{bmatrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 2 \\ 2 & 5 & 9 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{bmatrix} = -7$$
 captures vertical lines features

c)
$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & 1 \\ -1 & -1 & 1 \end{bmatrix}$$

captures diagonal lines features

4)

4) pooling using stride of /

$$\begin{bmatrix} 1 & 1 & 2 \\ 2 & 5 & 9 \\ 1 & 2 & 0 \end{bmatrix} = \text{max is } 9$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 5 & 9 & 7 \\ 2 & 0 & 1 \end{bmatrix} \text{ max is } 9$$

⋮ and so on took a 6 by 6
⋮ to a 4 by 4 image

result

$$\begin{bmatrix} 9 & 9 & 9 & 7 \\ 9 & 9 & 9 & 7 \\ 6 & 4 & 7 & 7 \\ 6 & 4 & 7 & 7 \end{bmatrix}$$

5)

- $3 \times 4 \times 4 = 48$ weights.
- After pooling, there's $3 \times 5 \times 5$, so 75 ReLU operations performed on the forward pass.
- $48 + 4 \times 75 = 348$ weights.
- It seems logical the answer is true. That another neural network with same architecture can represent the same classifier. Not sure how to explain in depth.
- It takes too long to train as there will be too many parameters with fully connected.

6) code

```

M reset_graph()

X = tf.placeholder(tf.float32, shape=(None, height, width, 1))
feature_maps = tf.constant(fmap)
# TODO :: Apply feature_maps to input image X with tf.nn.conv2d, the strides is 1 for each dimension of input. Expect 1
# line of code
convolution = tf.nn.conv2d(X, feature_maps, strides=[1,1,1,1], padding="SAME")

M # TODO :: create a session and evaluate the convolution operation
with tf.Session() as sess:
    output = convolution.eval(feed_dict={X: images})

# TODO :: Can you build two filters that can sharpen and blur an image respectively. Test your feature by applying
# it on the china.jpg. Expect 3 lines of code
fmap = np.zeros(shape=(7, 7, 1, 2), dtype=np.float32)
# Dr. Strachan showed this in live session notes
fmap[3, 3, 0, 0] = 1
fmap[:, :, 0, 1] = 1/49

plot_image(fmap[:, :, 0, 0])
plt.show()
plot_image(fmap[:, :, 0, 1])
plt.show()

```

```

reset_graph()

with tf.name_scope("inputs"):
    X = tf.placeholder(tf.float32, shape=[None, n_inputs], name="X")
    X_reshaped = tf.reshape(X, shape=[-1, height, width, channels])
    y = tf.placeholder(tf.int32, shape=[None], name="y")

# TODO : build two convolution layers here, use the hyperparameters we defined earlier, use relu as activation function
# name them as "conv1" and "conv2" respectively. Expect 2 lines of code.
conv1 = tf.layers.conv2d(X_reshaped, filters=conv1_fmmaps, kernel_size=conv1_ksize, strides=conv1_stride, padding=conv1_pad,
    activation=tf.nn.relu, name="conv1")
conv2 = tf.layers.conv2d(conv1, filters=conv2_fmmaps, kernel_size=conv2_ksize, strides=conv2_stride, padding=conv2_pad,
    activation=tf.nn.relu, name="conv2")

with tf.name_scope("pool3"):
    # TODO : define max pooling layer here, the kernel size is 2 * 2, and the strides is 2 for both vertical and horizontal,
    # Use "VALID" as padding method (remember to flatten the pooling result), expect 2 lines of code.
    pool3 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="VALID")
    pool3 = tf.reshape(pool3, shape=[-1, pool3_fmmaps * 49])

with tf.name_scope("fc1"):
    # TODO : define a fully connected layer for flattened pooling result. Name it as "fc1", expect 1 line of code.
    fc1 = tf.layers.dense(pool3, n_fc1, activation=tf.nn.relu, name="fc1")

with tf.name_scope("output"):
    # TODO : define the output layer and the result goes into softmax layer to get the prediction probability.
    # name the output layer as "output", name the softmax layer as "Y_proba", expect 2 lines of code.
    output = tf.layers.dense(fc1, n_outputs, name="output")
    Y_proba = tf.nn.softmax(output, name="Y_proba")

with tf.name_scope("train"):
    # TODO : define a loss function with sparse_softmax_cross_entropy_with_logits, and optimize the loss with AdamOptimizer
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=output, labels=y)
    loss = tf.reduce_mean(xentropy)
    optimizer = tf.train.AdamOptimizer()
    training_op = optimizer.minimize(loss)

with tf.name_scope("eval"):
    # TODO : define the accuracy of the model with tf.reduce_mean, you may want to use tf.nn.in_top_k, expect 2 lines of the
    correct = tf.nn.in_top_k(output, y, 1)
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

with tf.name_scope("init_and_save"):
    init = tf.global_variables_initializer()
    saver = tf.train.Saver()

```